

## INDEX

NAME: \_\_\_\_\_

REG NO: \_\_\_\_\_

Ex.No :

Date:

## Data Pre-processing & EDA using Pandas and Matplotlib.

**AIM:**

**ALGORITHM:**

**PROGRAM**

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
# Step 1: Given dataset (with missing values added)
```

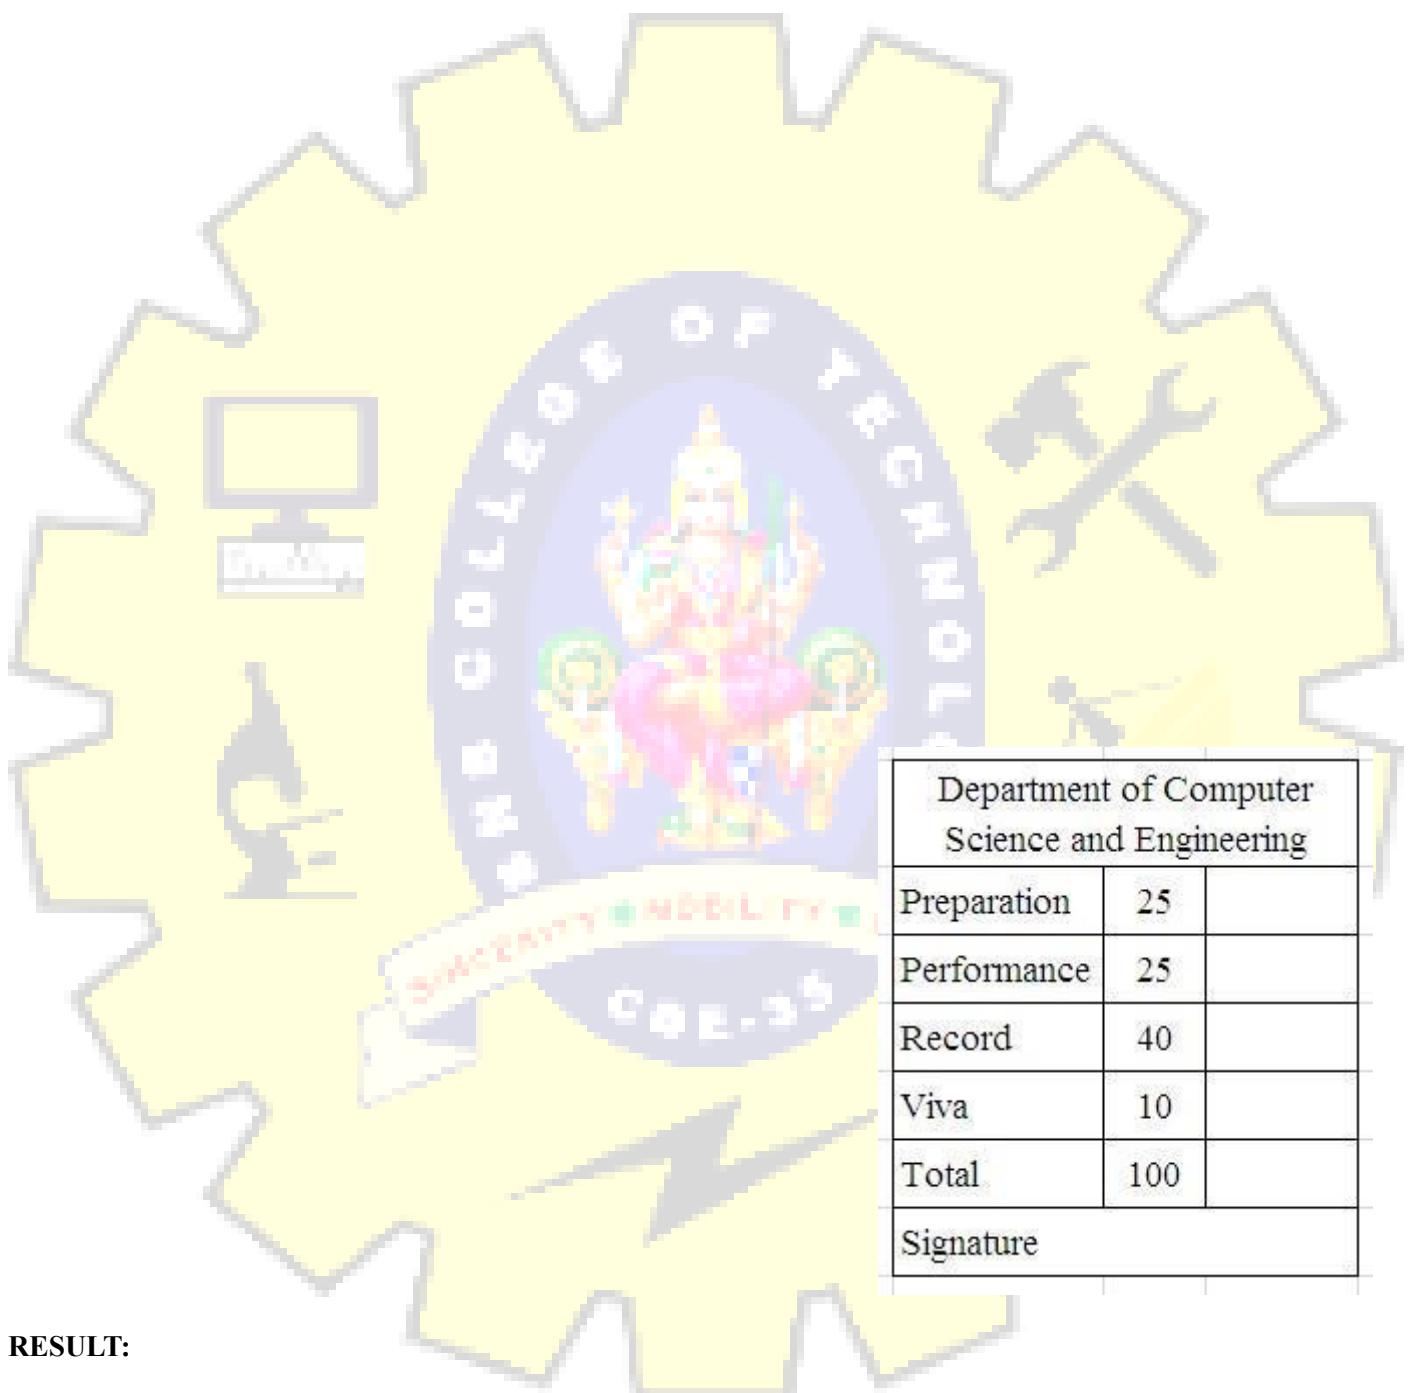
```
data = {
```

```
'Name': ['Jai','Princi','Ganga','Anuj','Ravi','Natasha','Riya'],  
'Age': ['17','18',np.nan,'19','18','17',np.nan],  
'Gender': ['M','F','F','M','M','F','F'],  
'Marks': ['90','76','86',np.nan,'82','88','67']  
}  
  
df = pd.DataFrame(data)  
print("Original Dataset:\n")  
print(df)  
  
# Step 2: Data type conversion  
df['Age'] = pd.to_numeric(df['Age'])  
df['Marks'] = pd.to_numeric(df['Marks'])  
  
# Step 3: Handling missing values  
df['Age'].fillna(df['Age'].mean(), inplace=True)  
df['Marks'].fillna(df['Marks'].mean(), inplace=True)  
print("\nAfter Handling Missing Values:\n")  
print(df)  
  
# Step 4: Encoding categorical variables  
df_encoded = pd.get_dummies(df, columns=['Gender'], drop_first=True)  
print("\nAfter Encoding Categorical Variables:\n")  
print(df_encoded)  
  
# Step 5: Normalization (Min-Max Scaling)  
df_encoded['Age'] = (df_encoded['Age'] - df_encoded['Age'].min()) / \  
(df_encoded['Age'].max() - df_encoded['Age'].min())  
df_encoded['Marks'] = (df_encoded['Marks'] - df_encoded['Marks'].min()) / \  
(df_encoded['Marks'].max() - df_encoded['Marks'].min())  
print("\nAfter Normalization:\n")  
print(df_encoded)  
  
# Step 6: Data Visualization
```

```
# Bar chart for Marks  
plt.bar(df['Name'], df['Marks'])  
plt.xlabel("Student Name")  
plt.ylabel("Marks")  
plt.title("Marks of Students")  
plt.show()  
  
# Pie chart for Gender distribution  
df['Gender'].value_counts().plot.pie(autopct='%1.1f%%')  
plt.title("Gender Distribution")  
plt.ylabel("")  
plt.show()
```

OUTPUT





**RESULT:**

NAME: \_\_\_\_\_

REG NO: \_\_\_\_\_

Ex.No :

Date:

## Linear Regression and Model Evaluation using MAE, MSE, R<sup>2</sup>, and plot residuals

**AIM:**

**ALGORITHM:**

**PROGRAM**

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Step 1: Create dataset (with error)
data = {
    'Hours_Studied': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Marks_Scored': [38, 41, 50, 52, 58, 63, 65, 71, 76, 82] # noisy data
}
df = pd.DataFrame(data)
print(df)

# Step 2: Define independent and dependent variables
X = df[['Hours_Studied']]
y = df['Marks_Scored']

# Step 3: Split dataset
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=1)

# Step 4: Train Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Step 5: Predict
y_pred = model.predict(X_test)

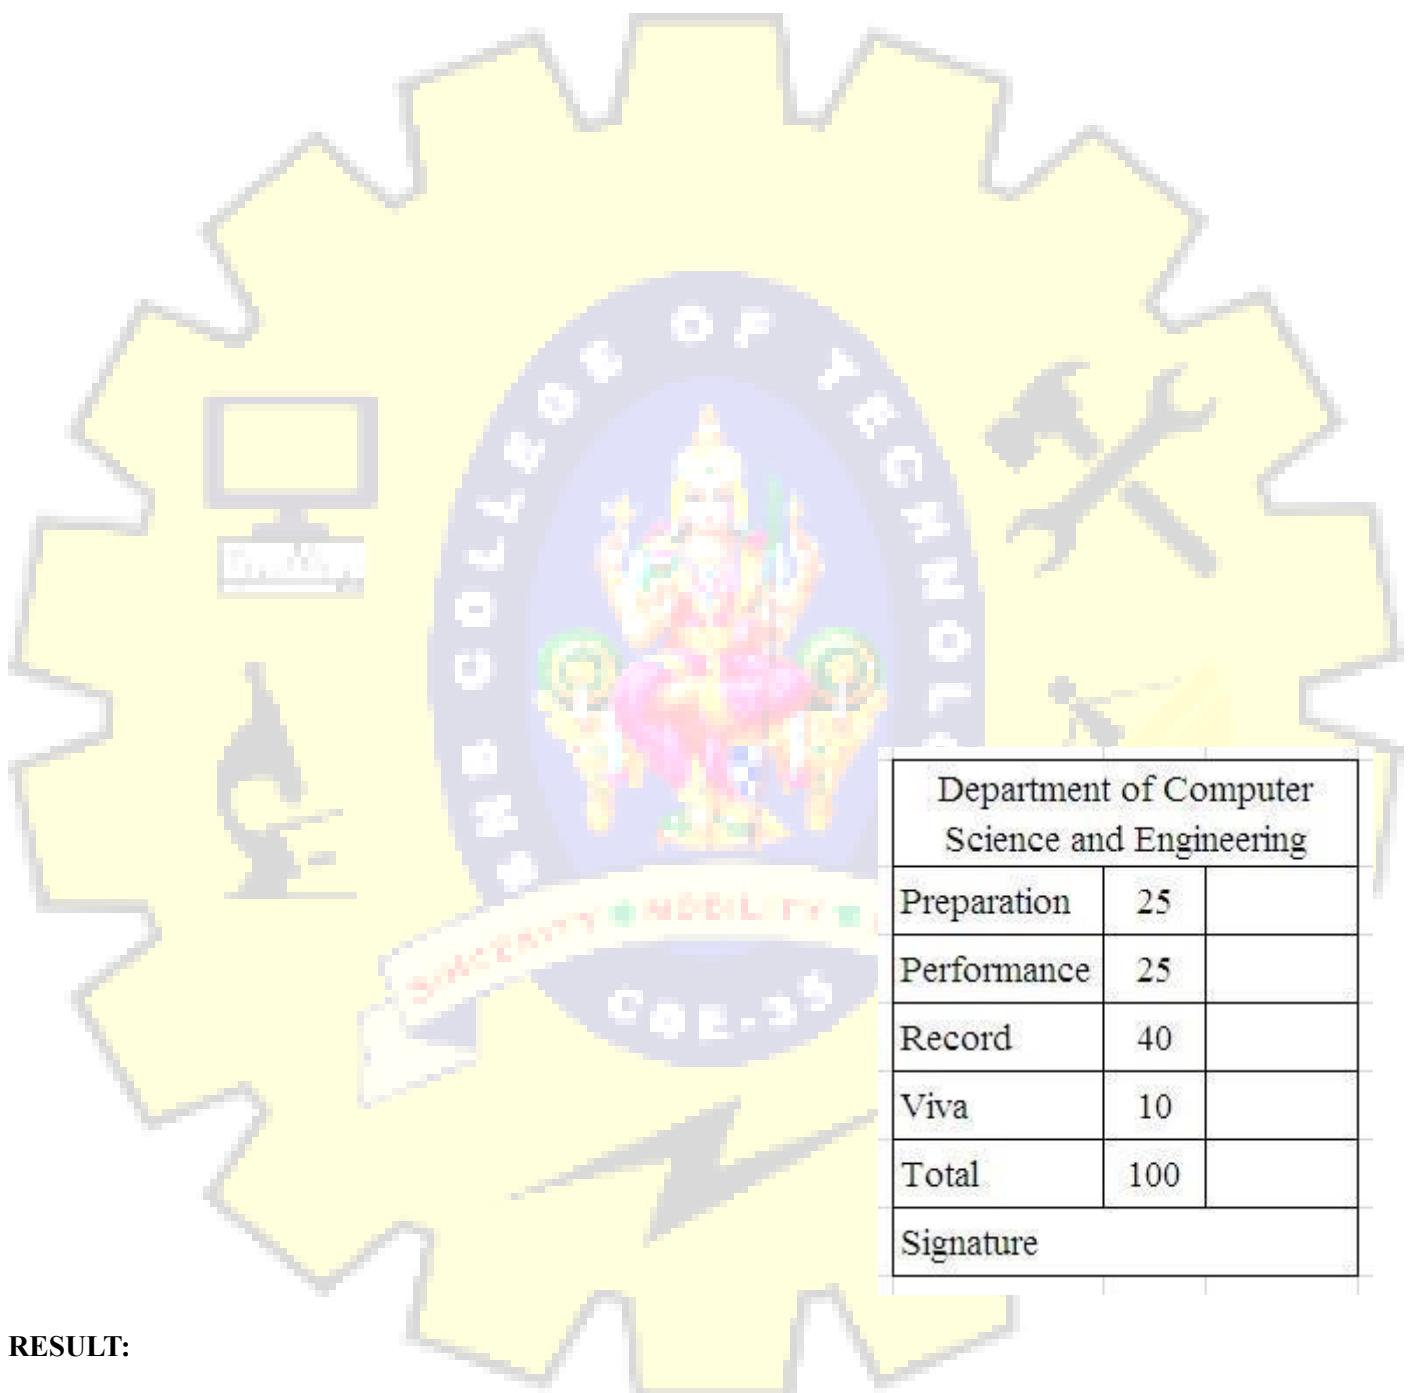
# Step 6: Evaluation metrics
print("\nEvaluation Metrics:")
print("MAE:", mean_absolute_error(y_test, y_pred))
print("MSE:", mean_squared_error(y_test, y_pred))
print("R2 :", r2_score(y_test, y_pred))

# Step 7: Residual plot
residuals = y_test - y_pred
plt.scatter(y_pred, residuals)
```

```
plt.axhline(0)
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residual Plot (With Error Data)")
plt.show()
```

**OUTPUT**





Department of Computer Science and Engineering		
Preparation	25	
Performance	25	
Record	40	
Viva	10	
Total	100	
Signature		

**RESULT:**

NAME: \_\_\_\_\_

REG NO: \_\_\_\_\_

Ex.No :

Date:

## Logistic Regression for Classification

**AIM:**

**ALGORITHM:**

**PROGRAM**

```
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import confusion_matrix, accuracy_score  
from sklearn.metrics import roc_curve, auc, ConfusionMatrixDisplay
```

```
#Step 1: Dataset
X = np.array([29, 15, 33, 28, 39]).reshape(-1, 1) # Hours studied
y = np.array([0, 0, 1, 1, 1]) # Pass(1)/Fail(0)

#Step 2: Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Step 3: Train Logistic Regression Model
model = LogisticRegression()
model.fit(X_train, y_train)

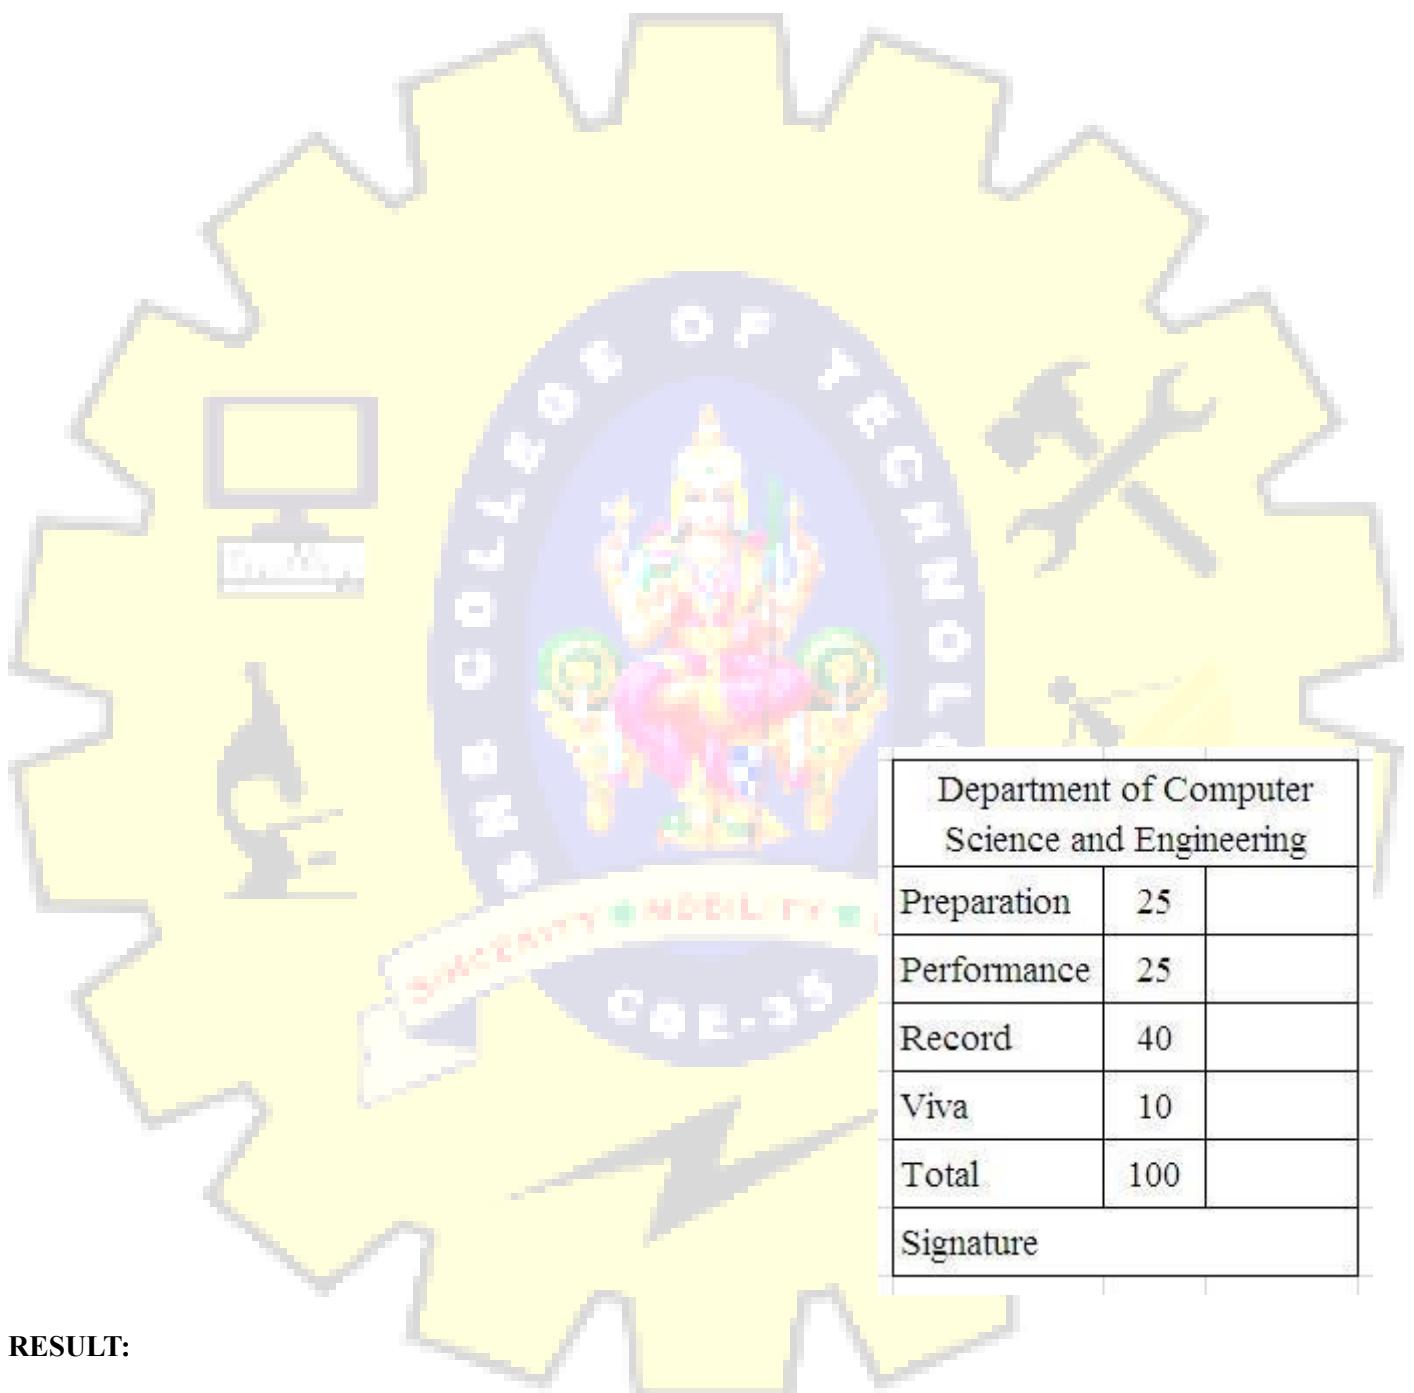
#Step 4: Predictions
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1]

#Step 5: Confusion Matrix & Accuracy
m = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=m)
disp.plot()
plt.title("Confusion Matrix")
plt.show()
print("Accuracy:", accuracy_score(y_test, y_pred))

#Step 6: ROC Curve & AUC
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label="ROC Curve (AUC = %.2f)" % roc_auc)
plt.plot([0, 1], [0, 1], linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
```

```
plt.legend()  
plt.show()  
  
#Step 7: Probability for Given Hours (33)  
prob_33 = model.predict_proba([[33]])[0][1]  
print("Probability of passing with 33 hours study:", prob_33)  
  
# Step 8: Minimum Hours for >95% Probability  
for hours in range(0, 61):  
    prob = model.predict_proba([[hours]])[0][1]  
    if prob > 0.95:  
        print("Minimum hours required for >95% pass probability:", hours)  
        print("Probability at this hour:", prob)  
        break
```

OUTPUT



Department of Computer  
Science and Engineering

Preparation	25	
Performance	25	
Record	40	
Viva	10	
Total	100	
Signature		

**RESULT:**

NAME: \_\_\_\_\_

REG NO: \_\_\_\_\_

Ex.No :

Date:

## K-Nearest Neighbors (K-NN) Classifier

**AIM:**

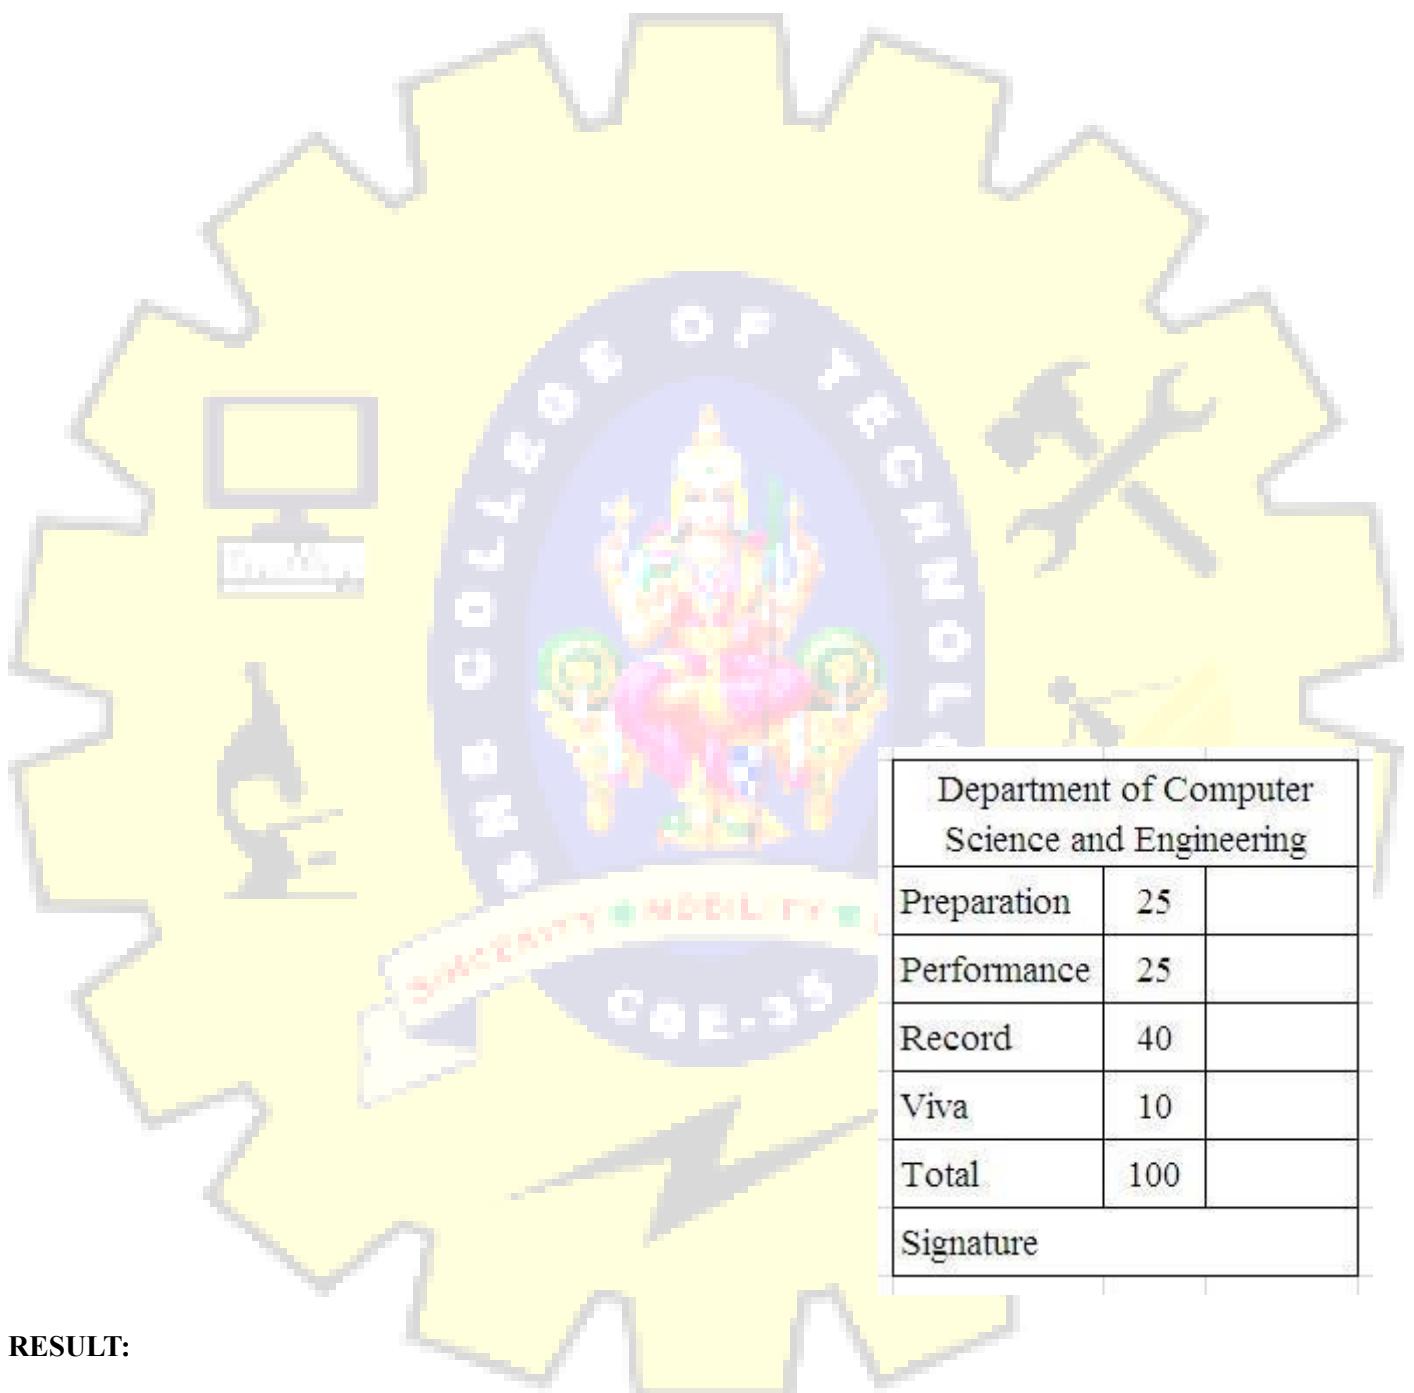
**ALGORITHM:**

**PROGRAM**

```
import numpy as np  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.preprocessing import StandardScaler  
# -----  
# Step 1: Dataset (Height, Weight)  
# -----
```

```
X = np.array([[167, 51],[182, 62],[176, 69],[173, 64],[172, 65],[174, 56],[169, 58],  
[173, 57],[170, 55]])  
  
y = np.array(["Underweight","Normal","Normal","Normal","Normal","Underweight",  
"Normal","Normal","Normal"])  
  
# New data point to classify  
  
new_point = np.array([170, 57])  
  
# Step 2: Normalization  
  
scaler = StandardScaler()  
  
X_scaled = scaler.fit_transform(X)  
  
new_point_scaled = scaler.transform(new_point)  
  
# Step 3: Apply k-NN (Euclidean Distance)  
  
k = 3 # You can change k value  
  
model = KNeighborsClassifier(n_neighbors=k, metric='euclidean')  
  
model.fit(X_scaled, y)  
  
# Prediction  
  
prediction = model.predict(new_point_scaled)  
  
print("Predicted Class for (170 cm, 57 kg):", prediction[0])
```

OUTPUT



Department of Computer Science and Engineering		
Preparation	25	
Performance	25	
Record	40	
Viva	10	
Total	100	
Signature		

**RESULT:**